Transformers for Time Series Forecasting

Marco Zanotti

University Milano-Bicocca

Contents

1. The TSF Problem

Time series forecasting (TSF) is the task of predicting future values of a given sequence based on previously observed values.

The TSF problem may be essentially identified by the following aspects:

▶ **Prediction objective**: point forecasting vs probabilistic forecasting

▶ **Forecast horizon**: short-term vs long-term forecasting

▶ **Input-Output dimension**: univariate vs multivariate forecasting

▶ **Forecasting task**: single-step vs multi-step forecasting

The TSF problem is usually faced with statistical models (ARIMA, ETS) or deep learning models (RNN, LSTM).

The main challenges of the TSF problem are:

▶ **uncertainty** increases as the forecast horizon increases

▶ difficulty in capturing **multiple complex patterns** over time

▶ difficulty in capturing **long-term dependencies** (critical for long-term forecasting)

▶ difficulty to handle **long input sequences**

2. Vanilla Transformer

▶ Based on Encoder-Decoder
architecture

▶ Each encoder and decoder layer is
composed of an attention layer and
a feed-forward layer

▶ Uses self-attention mechanism to
access any part of the sequence
history

▶ Positional encoding allows to
account for element positions
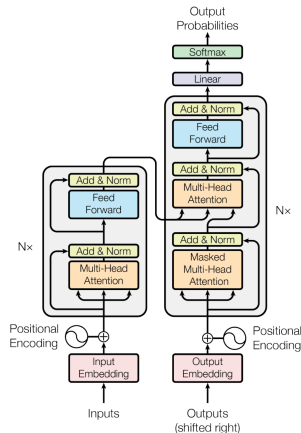
Can vanilla Transformers be used for TSF?

The TSF problem can be seen as a sequence learning problem such as machine translation.

Main ingredients allowing to use vanilla Transformers for TSF:

▶ **Multi-head Self-attention** mechanism allows to access any part of the sequence history, capturing both short-term and long-term dependencies (but it is invariant to the order of elements in a sequence)

▶ **Positional encoding** allows to account for the sequence ordering

▶ **Masked self-attention** allows to avoid information leakage from future

Can vanilla Transformers be used for TSF?

Just few changes are needed to adapt Transformers to TSF:

▶ **Remove the final activation** function (softmax) from the output layer and set the dimension of the linear layer equal to the forecasting horizon

▶ **Adapt the structure** to the desired forecasting task (single-step or multi-step)

Problems with vanilla Transformers

▶ **Positional encoding**: only the order in which two elements occur is taken into account, but their temporal distance is not

▶ **Computational complexity**: given a sequence of length L, the time and memory burden is $O(L^2)$, making it difficult to learn patterns in long time series

▶ **Simple Architecture**: the architecture does not include any component of typical importance in TSF (e.g. autocorrelation, decomposition, recurrent layers, etc.)

3. TSF Transformers

Transformers are **very appealing for long-term TSF** due to their ability to learn long-range dependencies.

Many solutions have been proposed to adapt Transformers to TSF, mainly in the direction to adopt **more efficient attention** and **expand the architecture**.


(a) Preprocessing | (b) Embedding | (c) Encoder | (d) Decoder
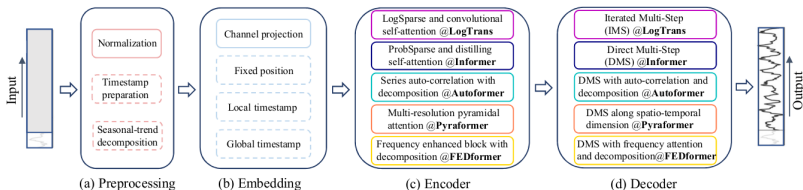
Informer

Informer employs two major improvements:

▶ **ProbSparse Attention Mechanism**, which replaces the standard self-attention used in the vanilla transformer with a sparse attention so to achieve $O(LlogL)$ complexity

▶ **Distilling Module**, which reduces the input size between encoder layers into its half slice, removing redundancy and reducing the computational burden

Informer - Architecture
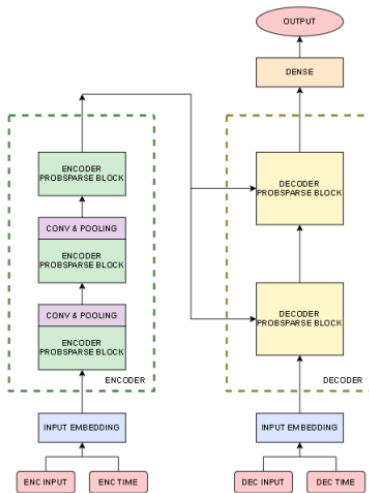
Informer - ProbSparse Attention

The main idea of **ProbSparse** is that just a small subset of queries ("active") effectively contribute to the attention mechanism. The ProbSparse attention selects the "active" queries, and creates a reduced query matrix, which is then used to calculate the attention weights, reaching $O(LlogL)$ complexity.

(a) Full Attention   (b) Sparse Attention

Informer - Convolution Layers

Because of the ProbSparse self-attention, the encoder's feature map has some redundancy that can be removed. The **distilling** operation is used to reduce the input size between encoder layers.

In practice, Informer's distilling operation just adds **1D convolution layers**, along the time dimension, with max pooling between each of the encoder layers.

$$X_{n+1} = MaxPool(\,ELU(\,Conv1d(Xn)\,)\,)$$

This reduces by half the size of the data between the feature space, improving both memory and computational efficiency.

Autoformer

Autoformer builds upon two traditional time series analysis methods:

▶ **Decomposition Layer**, which allows to decompose the time series into seasonality and trend-cycle components, enhancing the model's ability to capture these components accurately

▶ **Autocorrelation Attention Mechanism**, which replaces the standard self-attention used in the vanilla transformer with an autocorrelation mechanism, allowing to capture the temporal dependencies in the frequency domain and achieving $O(LlogL)$ complexity

Autoformer - Architecture
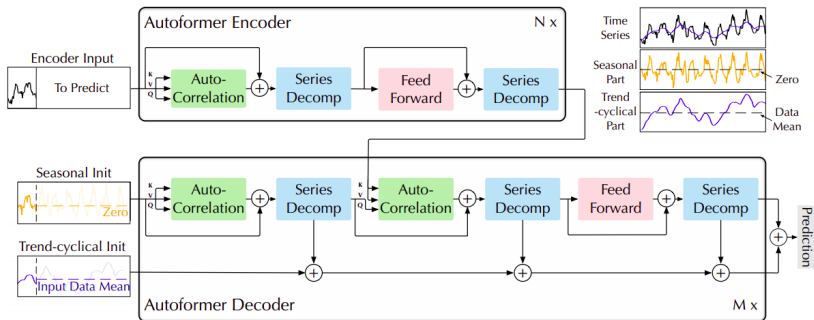
Autoformer - Decomposition Layer

Autoformer incorporates **decomposition blocks** as an inner operation of the model.

Encoder and decoder use decomposition blocks to extract and aggregate the trend-cyclical and seasonal components from the series progressively, so to make raw data easier to predict.

For an input series X_t with length L, the decomposition layer returns X_{trend} and $X_{seasonal}$, both of length L. In practice, X_{trend} is extracted using some form of moving-average and $X_{seasonal}$ is then obtained by difference.

Autoformer - Autocorrelation Attention

Autoformer uses **autocorrelation within the self-attention** layer, extracting frequency-based dependencies from (Q, K). The autocorrelation block measures the **time-delay similarity** and aggregates the top n similar sub-series to reduce complexity.

(a) Full Attention

(b) Auto-Correlation

4. Conclusions

Conclusions

▶ Transformers are very appealing for TSF due to their ability to learn long-range dependencies

▶ Many solutions have been proposed to adapt Transformers to TSF, mainly in the direction to adopt more efficient attention and expand the architecture

▶ Informer and Autoformer are two of the most successful solutions, and they both achieve $O(LlogL)$ complexity

▶ Empirical results show that Transformer models are able to reach SOTA performance in TSF, and are among the best methods for long-term forecasting

Bibliografy

Ailing Z., et al., 2023, 'Are Transformers Effective for Time Series Forecasting?', AAAI

Haixu W., et al., 2021, 'Autoformer: Decomposition transformers with auto-correlation for long-term series forecasting', NeurIPS

Haoyi Z., et al., 2021, 'Informer: Beyond efficient transformer for long sequence time-series forecasting', AAAI

Lara-Benitez P., et al., 2021, 'Evaluation of the Transformer Architecture for Univariate Time Series Forecasting', Advances in Artificial Intelligence, CAEPIA

Qingsong W., et al., 2022, 'Transformers in Time Series: A Survey', AAAI

Thank you!

Appendix

Autocorrelation

In theory, given a time lag τ, **autocorrelation** for a single discrete variable Y is used to measure the "relationship" between the variable at time t to its past value at time $t - \tau$.

$$Autocorrelation(\tau) = \mathcal{R}(\tau) = Corr(Y_t, Y_{t-\tau})$$
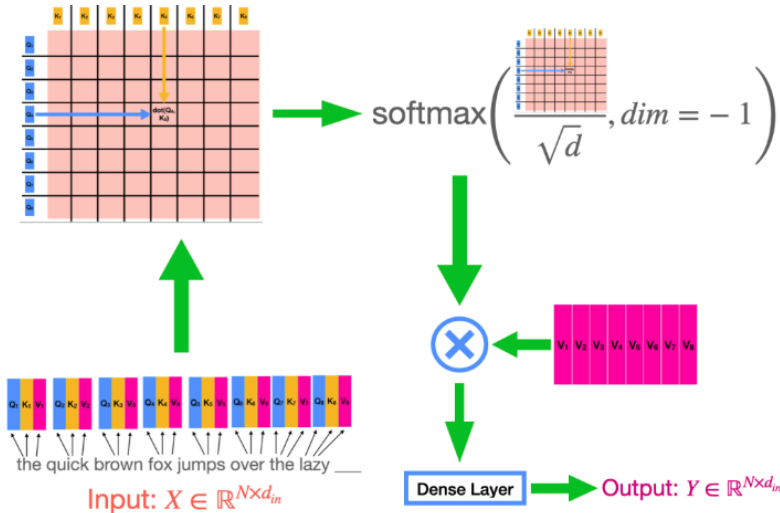
Time Series Decomposition

In time series analysis, **decomposition** is a method of breaking down a time series into three systematic components: **trend-cycle**, **seasonal variations**, and **random fluctuations**.
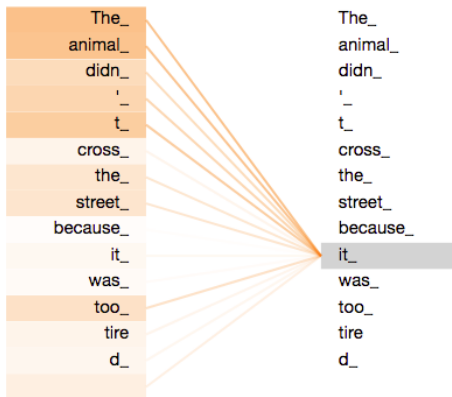
Attention Mechanism

$$\text{softmax}\left(\frac{}{\sqrt{d}}, dim = -1\right)$$

Input: $X \in \mathbb{R}^{N \times d_{in}}$

Output: $Y \in \mathbb{R}^{N \times d_{in}}$

Attention Mechanism

Machine translation example:
"The animal didn't cross the street because it was too tired"

Informer - ProbSparse Attention

$$Attention(Q, K, V) = softmax(\frac{Q_{reduced}K^T}{\sqrt{d_k}})V$$

$$M(q_i, K) = \max_j \frac{q_i K_j^T}{\sqrt{d_k}} - \frac{1}{L_k} \sum_{j=1}^{L_k} \frac{q_i K_j^T}{\sqrt{d_k}}$$

Queries (50) & Keys (50) Feature Map

The "Active" Query in Self-Attention

The "Lazy" Query in Self-Attention

Autoformer - Autocorrelation Attention

$$\tau_1, ..., \tau_k = arg\ Top - k\ (\mathcal{R}_{Q,K}(\tau))$$

$$Attention(Q, K, V) = \sum_{i=1}^{k} Roll(V, \tau_i) \times softmax(\mathcal{R}_{Q,K}(\tau_i))$$

RNN & LSTM vs Transformers

RNN:

▶ process information sequentially
▶ long-term memory loss
▶ vanishing gradients problem
▶ computation issues due to sequential nature

LSTM:

▶ improved the long-term memory loss
▶ solved the vanishing gradients problem
▶ still computational issues due to sequential nature